

The background of the entire image is a dark blue field filled with a pattern of red dots. These dots are arranged in a way that they form a large, faint circular shape in the center, with the density of the dots increasing towards the perimeter of this circle. The dots are of uniform size and color.

HUST

ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

ONE LOVE. ONE FUTURE.



ĐẠI HỌC
BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY
OF SCIENCE AND TECHNOLOGY

Multiscale and Multilayer Contrastive Learning for Domain Generalization

Group 26 members

Vu Thuong Tin - 20230091

Chu Anh Duc - 20230081

Nguyen Xuan Khai - 20235508

ONE LOVE. ONE FUTURE.

1. Introduction
2. The $M^2 - CL$ Architecture
3. Experimental
4. Qualitative Visualizations
5. Conclusion
6. Appendix

Vu Thuong Tin

Baseline & Problem Setting

- ▶ DG problem structure
- ▶ Train/test pipeline
- ▶ ERM baseline
- ▶ Dataset loading & domain split
- ▶ Why ERM fails

Chu Anh Duc

Core Model

- ▶ M^2 architecture
- ▶ Extraction block
- ▶ Spatial Dropout
- ▶ Multi-scale pooling
- ▶ Projection head

Nguyen Xuan Khai

Loss & Experiments

- ▶ Contrastive loss
- ▶ Hyperparameters (α, τ)
- ▶ Evaluation & accuracy
- ▶ t-SNE visualization
- ▶ Saliency maps

The Core Objective

Train a parametric model $f(\cdot; \theta)$ on a set of N available **source domains** $\mathcal{S} = \{S_1, S_2, \dots, S_N\}$ such that it generalizes robustly to K completely **unseen target domains** $\mathcal{T} = \{T_1, T_2, \dots, T_K\}$.

Why Is This Hard?

- ▶ **Zero access** to target-domain data during training no direct supervision signal.
- ▶ Models must learn **domain-invariant** features from source data alone.
- ▶ Distribution shift at test time causes catastrophic accuracy drops in naive ERM.

Representative DG Approaches

- ▶ **ERM:** Aggregates all source data learns domain-specific shortcuts.
- ▶ **CORAL / MMD:** Aligns first/second-order statistics **single layer only**, loses geometric detail.
- ▶ **RSC:** Drops dominant features at training ignores multi-scale structure.
- ▶ **Mixup:** Interpolates images across domains does not enforce feature invariance.

What M^2 -CL Contributes

- ▶ **Multi-scale & multi-layer** extraction captures geometry at every depth.
- ▶ **Spatial Dropout** in each block reduces reliance on style shortcuts.
- ▶ **Layer-wise contrastive loss** encourages same-class alignment across domains at *all* scales simultaneously.

Why ERM Fails: Shortcut Learning

ERM Trap

- ▶ Learns **spurious cues** instead of structure.
- ▶ Fails under domain shift.

M^2 -CL Motivation

- ▶ Enforce domain-invariant features.
- ▶ Align same-class features at multiple scales.




Fig: ERM relies on “Grass” (shortcut) instead of actual “Cow” geometry.


The Last-Layer Problem

- ▶ Standard CNNs use **deepest features** only.

The Invariant Signal

- ▶ **Early/Mid Layers:** Hold shapes and edges.

⇒ Corrupted by style and filters out raw structure details.




Causal vs. Style Attribute Disentanglement

Latent Space Split: $X \approx (a, d)$

- ▶ **Causal (a):** Class-specific geometry (Invariant).
- ▶ **Style (d):** Domain-specific context and texture (Noise).

The DG Goal


Isolate **Causal** a from **Style** d to ensure the model only relies on invariant geometry for prediction.




The M^2 Architectural Paradigm

Framework Concept

Leverages parallel multi-scale concentration blocks mapped across multiple intermediate layers of a standard backbone network.





Concentration pipeline layout.

Internal Sequence Mechanics


1. **1×1 Convolutional Layer:** Compresses high-dimensional channels c down to $\lfloor c/r \rfloor$ using a reduction factor r .
2. **Spatial Dropout Layer:** Mitigates pixel correlation leakage.
3. **Multi-Scale Max Pooling:** Captures spatial features before reduction.
4. **MLP Projection Vector:** Maps pooled features into vectors via MLPs.

Vanilla Dropout

Drops **individual pixels** - adjacent pixels remain correlated; style information **leaks through**.

Spatial Dropout

Drops **entire feature maps** ($H \times W$ channels). Encourages the network to rely on alternative feature maps, **reducing style shortcut reliance** at the source.



Geometric Constraints

- ▶ **Pull:** Same class, cross domains.
- ▶ **Push:** Different classes, same style.

Alignment Measure ($\|u\| = 1$)

$$p^{(l)}(c) = \frac{\sum_{i,j \in N_c} \exp(u_i^T u_j / \tau)}{\sum_{i,a \in N_b} \exp(u_i^T u_a / \tau)}$$

Total Objective

$$\mathcal{L} = \mathcal{L}_{CE} + \alpha \sum_{l=1}^M \mathcal{L}^{(l)}$$

$$\mathcal{L}^{(l)} = - \sum_{c=1}^C \log p^{(l)}(c)$$

Optimizer

- ▶ **SGD, 30 epochs:** Stable convergence; training curves plateau before epoch 30.
- ▶ **LR = 0.001:** Standard for fine-tuning pretrained ResNet; avoids over-shooting.
- ▶ **Batch = 128:** Contrastive loss requires enough positive/negative pairs per step.

Architecture & Loss

- ▶ $r \in \{2, 4, 6\}$: Compression ratio of the projection module.
- ▶ $\alpha \in [0.001, 0.1]$: Weight of the contrastive loss.
- ▶ $\tau \in [0.8, 1.4]$: Temperature in contrastive learning.

VLCS

10,729 images | 5 classes

Domains: PASCAL, LabelMe, Caltech, SUN



(a) VLCS

PACS

9,991 images | 7 classes

Domains: Photo, Art, Cartoon, Sketch



(b) PACS

Office-Home

15,588 images | 65 classes


Domains: Art, Clipart, Product, Real-World



(c) Office-Home

Training Convergence on PACS (ResNet-18)

Training Convergence on PACS (ResNet-18)



Results: PACS and VLCS (ResNet-18)

Top-1 accuracy on PACS, VLCS (ResNet-18)

Method	Art	Cartoon	Photo	Sketch	PACS Avg	Caltech	Labelme	Sun	Voc	VLCS Avg
ERM	76.53	74.80	94.95	67.89	78.54	96.87	65.03	71.62	71.60	76.28
RSC	79.26	76.12	95.07	70.88	80.34	97.68	66.33	73.53	72.75	77.57
CORAL	78.55	72.74	92.24	73.93	79.36	95.16	60.73	70.66	65.00	72.89
MIXUP	79.79	72.81	93.11	63.97	77.42	93.44	60.25	74.36	69.87	74.48
MMD	77.75	73.35	92.34	69.74	78.29	97.28	60.62	70.48	65.21	73.40
SagNet	78.71	74.90	94.79	70.48	79.72	95.76	61.22	69.61	71.90	74.62
SelfReg	79.46	76.32	95.93	71.60	80.83	92.53	62.56	70.27	72.15	74.38
ARM	79.77	75.33	94.43	<u>74.76</u>	81.07	95.86	60.19	70.66	72.83	74.89
EQRM	78.61	72.38	92.48	<u>71.13</u>	78.65	96.17	60.57	70.44	64.66	72.96
SAGM	78.22	75.14	94.79	70.46	79.65	96.37	66.22	71.14	71.35	76.27
M2	<u>80.64</u>	<u>76.91</u>	95.50	74.20	<u>81.81</u>	<u>98.24</u>	<u>67.02</u>	73.80	<u>72.92</u>	<u>78.00</u>
M2-CL	81.04	77.48	<u>95.70</u>	75.81	82.51	98.69	67.48	<u>74.05</u>	73.76	78.50

Dataset Table : Office-Home (ResNet-18)

Top-1 accuracy on Office-Home, ResNet-18.

Method	Art	Clipart	Product	Real-World	Avg
ERM	56.12	47.88	<u>72.47</u>	73.58	62.51
RSC	56.00	46.09	70.85	72.32	61.31
CORAL	55.46	49.39	71.75	73.28	62.47
MIXUP	56.37	<u>49.81</u>	70.87	73.61	62.66
MMD	55.67	48.52	71.68	73.19	62.27
SagNet	55.75	47.74	71.71	73.17	62.09
SelfReg	57.56	49.69	72.43	73.72	63.35
ARM	54.47	48.45	71.14	<u>73.79</u>	61.96
EQRM	56.08	48.34	71.57	72.94	62.23
SAGM	55.79	46.51	71.16	73.15	61.65
M2	<u>58.04</u>	49.30	73.04	73.50	<u>63.47</u>
M2-CL	58.77	50.86	72.20	74.91	64.18

Average Accuracy

- ▶ **ResNet-18:** M^2 -CL reaches **82.51** on PACS, **78.50** on VLCS, and **64.18** on Office-Home.
- ▶ The average score improves over the strongest baseline in each benchmark summary.

Interpretation

- ▶ M^2 improves feature representations through multi-scale reuse.
- ▶ M^2 -CL aligns same-class features across domains.
- ▶ Ablation studies identify the sources of these gains.

Ablation Study on PACS and VLCS

Ablation study on PACS and VLCS evaluating the effects of the projection type (pipe), compression ratio (r), dropout, and contrastive loss. The best and second-best results are highlighted in bold and underlined, respectively.


pipe	r	drop	loss	A	C	P	S	PACS Avg	C	L	S	V	VLCS Avg
c	2	-	-	72.34	72.29	94.19	69.51	78.23	93.87	<u>60.91</u>	65.89	61.91	71.22
c	4	-	-	74.38	72.74	90.13	69.71	77.21	94.14	<u>59.01</u>	67.82	65.42	69.62
c	6	-	-	75.72	70.12	94.26	73.33	77.92	91.93	56.81	67.28	64.10	68.95
c	2	yes	-	75.67	73.29	92.16	<u>75.95</u>	77.93	93.73	57.72	66.34	62.50	71.74
c	4	yes	-	74.36	74.05	92.57	76.45	79.37	<u>96.07</u>	56.28	67.07	64.30	65.81
c	6	yes	-	75.97	72.77	93.01	74.73	76.81	<u>92.94</u>	59.73	66.15	64.51	70.09
p	2	-	-	76.09	55.88	68.35	60.92	64.40	95.60	56.97	66.25	73.15	<u>73.96</u>
p	4	-	-	74.72	74.62	92.01	70.32	77.64	95.20	58.10	66.20	64.90	70.29
p	6	-	-	75.38	69.99	91.21	69.47	77.77	95.23	56.81	69.13	65.34	70.55
p	2	yes	-	74.61	72.39	87.08	73.95	78.88	95.73	58.33	68.44	63.74	70.49
p	6	yes	-	<u>76.73</u>	73.98	93.41	73.56	<u>79.84</u>	93.19	56.05	66.62	64.24	70.43
p	4	yes	-	76.39	75.98	<u>95.58</u>	74.49	79.15	94.54	59.77	67.26	<u>74.55</u>	73.14
p	4	yes	yes	80.43	<u>75.43</u>	96.87	74.42	80.10	96.09	61.15	<u>69.05</u>	75.70	74.15

Table: Reference Insights (PACS / VLCS, ResNet-18)

- ▶ **Best full configuration:** $p, r = 4$, $drop = yes$, $loss = yes$ reaches **80.10** PACS Avg and **74.15** VLCS Avg.
- ▶ **Contrastive loss contribution:** adding the loss to the same $p, r = 4$, $drop = yes$ setting improves PACS from **79.15** to **80.10** and VLCS from **73.14** to **74.15**.
- ▶ **Spatial dropout effect:** for $p, r = 4$, dropout improves PACS from **77.64** to **79.15** and VLCS from **70.29** to **73.14**.

Feature Space: t-SNE Visualization

Feature Space Visualization — PACS (Sketch target domain, ResNet-18)



ERM (Left)

Same-class points from different domains are **scattered** - the model encodes domain identity, not class identity.

M²-CL (Right)

The layer-wise contrastive loss **reduces intra-class domain scatter**: samples from the same class become more compact, while inter-class boundaries remain clearer.

Insight: The visualization suggests more domain-robust representations and is consistent with the contrastive objective.

Saliency Maps: What the Model Focuses On

PREDICTION

Volcano


Confidence 100.0%

Volcano	100.0%
Torch	0.0%
Tent	0.0%
Yurt	0.0%
Alp	0.0%

Input Image




ERM (Baseline)



Scattered attention · 61.1% of pixels active

M²-CL (Ours)



Focused attention · 19.7% of pixels active

Bright blue pixels = high gradient magnitude · Dark = model ignores

Guided backprop (M²-CL) is 3.1x more spatially concentrated than vanilla gradient (ERM)

ERM

Saliency spreads across **backgrounds and textures** - the model fires on domain-specific context (grass, sky) rather than object structure.

M²-CL

Saliency tends to concentrate more on **object contours and edges** across Photo, Art, Cartoon, and Sketch domains.

Insight: The contrastive loss encourages the network to rely more on *structural geometry* and less on style-specific cues.

Core Achievements

- ▶ **Architecture:** Combines multi-scale, multi-layer intermediate features while reducing reliance on domain-specific styles.
- ▶ **Loss Optimization:** Introduces a layer-wise contrastive loss that encourages class-level alignment across extraction levels.
- ▶ **Empirical Validation:** Shows competitive improvements over strong domain generalization baselines across three selected benchmark datasets.

Current Computational Drawbacks

- ▶ **Memory Overhead:** Concatenating multiple intermediate feature maps creates dense hypercolumns before the classification head, increasing memory usage.
- ▶ **Batch Size Dependency:** Like most contrastive methods, M^2 -CL requires large training batches to construct effective positive/negative pairs.

Future Research Avenues

Explore cross-layer attention mechanisms, analyze object-level evidence more rigorously, and evaluate alternative similarity metrics like KL divergence to refine latent distribution shapes.


A decorative graphic on the left side of the slide. It features a dark blue background with a large, stylized circular pattern composed of many small red dots. The dots are arranged in concentric, slightly offset rings, creating a sense of depth and movement. The word "HUST" is centered within this pattern.

HUST

THANK YOU !



Method Comparison — All Datasets (ResNet-18)



Dataset Table: PACS and VLCS (ResNet-50)

Top-1 accuracy on PACS and VLCS, ResNet-50.

Method	Art	Cartoon	Photo	Sketch	Avg	Caltech	Labelme	Sun	Voc	Avg
ERM	83.54	78.47	96.51	73.64	83.04	97.01	62.22	71.79	76.09	76.78
RSC	82.08	80.26	96.51	75.45	83.57	97.81	64.44	74.60	75.53	78.10
CORAL	83.94	79.27	94.37	81.98	84.89	97.01	61.34	75.26	71.74	76.34
MIXUP	85.16	77.40	96.23	78.26	84.26	96.94	62.33	73.40	72.32	76.25
MMD	82.21	74.70	92.89	69.00	79.70	97.48	63.85	75.49	71.65	77.12
SagNet	85.23	78.47	93.20	77.71	83.65	94.14	60.43	70.49	69.96	73.75
SelfReg	83.33	79.06	93.16	77.24	83.20	92.49	60.73	66.80	73.58	73.40
ARM	82.14	79.10	92.97	73.02	81.81	96.01	63.98	66.26	71.79	74.51
EQRM	84.98	77.74	91.26	74.88	82.21	94.35	60.80	66.35	72.94	73.61
SAGM	81.35	78.99	93.24	79.39	83.24	97.39	64.23	69.30	76.17	76.77
M2	<u>86.02</u>	79.90	<u>97.03</u>	81.50	<u>86.11</u>	97.52	<u>64.77</u>	<u>76.02</u>	<u>76.72</u>	<u>78.76</u>
M2-CL	86.48	<u>80.05</u>	97.99	<u>81.70</u>	86.56	<u>97.60</u>	64.86	76.65	77.26	79.09

Dataset Table: Office-Home (ResNet-50)

Top-1 accuracy on Office-Home, ResNet-50.

Method	Art	Clipart	Product	Real-World	Avg
ERM	63.43	51.32	76.45	<u>78.13</u>	67.33
RSC	62.64	51.55	75.11	77.34	66.66
CORAL	61.27	<u>54.34</u>	73.35	77.83	66.70
MIXUP	63.23	51.91	<u>77.10</u>	76.39	67.16
MMD	59.63	49.37	75.28	74.71	64.75
SagNet	61.41	51.85	72.85	73.79	64.97
SelfReg	59.95	51.22	71.28	74.52	64.24
ARM	54.23	52.14	71.73	71.26	62.34
EQRM	63.52	52.32	73.28	76.45	66.39
SAGM	61.01	53.27	74.81	77.00	66.52
M2	<u>65.35</u>	53.90	76.80	78.74	<u>68.70</u>
M2-CL	66.46	54.82	78.24	77.85	69.34

Table: Sensitivity to Temperature τ

Sensitivity analysis for τ in M2-CL, ResNet-18.

τ	PACS	VLCS
0.01	78.10	76.15
0.1	80.03	77.27
0.2	81.37	74.71
0.4	82.07	72.03
0.6	83.70	74.15
0.8	84.67	75.21
1.0	80.07	76.24
1.2	81.34	75.11
1.4	83.14	75.69
1.6	82.17	76.84
1.8	80.51	74.66
2.0	78.72	74.99
10.0	79.92	74.04
100.0	81.84	73.74

Table: Sensitivity to Loss Weight α

Sensitivity analysis for α in M2-CL, ResNet-18.

α	PACS	VLCS
0.00	82.76	72.36
10e-5	80.41	73.56
10e-4	81.17	73.71
10e-3	82.29	74.38
10e-2	79.70	73.39
10e-1	64.04	55.96

Statistical Alignment

- ▶ **CORAL:** Aligns second-order statistics (covariance) across source domains.
- ▶ **MMD:** Minimizes Maximum Mean Discrepancy between domain distributions.

Data Augmentation / Interpolation

- ▶ **Mixup:** Interpolates images and labels across domain pairs.
- ▶ **SagNet:** Style-adversarial training to separate content from style.

Feature Regularization

- ▶ **RSC:** Randomly suppresses dominant activations to prevent shortcut reliance.
- ▶ **SelfReg:** Self-regularization via contrastive and distance-based losses.
- ▶ **ARM:** Context-conditional adaptation at inference time.
- ▶ **SAGM:** Sharpness-aware gradient matching for flat loss minima.
- ▶ **EQRM:** Empirical quantile risk minimization to reduce tail risk.

Hyperparameter Tuning

- ▶ **Temperature optimum:** PACS peaks at **84.67** with $\tau = 0.8$, while VLCS peaks at **77.27** with $\tau = 0.1$.
- ▶ **Default-scale behavior:** $\tau = 1.0$ remains competitive with **80.07** PACS and **76.24** VLCS.
- ▶ **Loss-weight sensitivity:** $\alpha = 10e^{-3}$ gives the best VLCS score, **74.38**, while $\alpha = 0$ gives the highest PACS score, **82.76**.

Alpha Bounds Limits

$$\alpha = 10e^{-1} \implies \text{PACS } 64.04, \quad \text{VLCS } 55.96$$